

Tracking, Retrieving, and Auditing for Coherent Epics in Online Narratives

Qiqi Duan*

The Hong Kong University of Science
and Technology (Guangzhou)
Guangzhou, China
qduan130@connect.hkust-gz.edu.cn

Chen Wang*

The Hong Kong University of Science
and Technology (Guangzhou)
Guangzhou, China
cwang018@connect.hkust-gz.edu.cn

Yuxiang Luo

The Hong Kong University of Science
and Technology (Guangzhou)
Guangzhou, China
yluo613@connect.hkust-gz.edu.cn

Nan Tang

The Hong Kong University of Science
and Technology (Guangzhou)
Guangzhou, China
nantang@hkust-gz.edu.cn

Leixian Shen

The Hong Kong University of Science
and Technology
Hong Kong, China
lshenaj@connect.ust.hk

Yuyu Luo[†]

The Hong Kong University of Science
and Technology (Guangzhou)
Guangzhou, China
yuyuluo@hkust-gz.edu.cn




Figure 1: The workflow of TRACE. When *writing*, (1) the author saves the initial text, and the system automatically extracts key entities to fill the Character Cards. (2) As the author adds more content, this knowledge base updates in real time. The author can keep writing in the editor without interruptions. When starting the *audit* phase, (3) the author clicks the “Audit” button to begin a consistency check. (4) The system analyzes the text against the knowledge base and highlights possible conflicts with inline notes and suggestions. In the *correct* phase, (5) the author has full control to accept, reject, or revise the text. (6) The author then continues writing, completing the “Write, Audit, Correct” loop, with the revised chapter merged back into the knowledge base.

*Both authors contributed equally to this research.

[†]Corresponding author



This work is licensed under a Creative Commons Attribution 4.0 International License.
CHI EA '26, Barcelona, Spain

© 2026 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-2281-3/26/04
<https://doi.org/10.1145/3772363.3798905>

Abstract

Long-form writing for online literature demands immense creative effort, yet authors are often bogged down by the high cognitive load of maintaining narrative consistency across millions of words. To address this, we started a 2-stage formative study with professional authors, deriving 3 design goals, motivating a “Write, Audit, Correct” workflow for author-controlled human-AI collaboration. Based on this, we present TRACE, a system leveraging dynamic

structured memory to check for narrative inconsistencies following the cognitive writing process. A user study with 12 authors provides early evidence that TRACE helps offload the cognitive burden of consistency checking, enabling them to remain focused and maintain their creative flow. Our work offers design implications for future AI tools that aim to augment, rather than automate human creativity.

CCS Concepts

• **Human-centered computing** → **Interactive systems and tools**; **User centered design**; • **Computing methodologies** → **Natural language generation**.

Keywords

Interactive Narrative, Human-AI Collaboration, Large Language Models, Creativity Support Tools.

ACM Reference Format:

Qiqi Duan, Chen Wang, Yuxiang Luo, Nan Tang, Leixian Shen, and Yuyu Luo. 2026. Tracking, Retrieving, and Auditing for Coherent Epics in Online Narratives. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 6 pages. <https://doi.org/10.1145/3772363.3798905>

1 Introduction

Online literature has evolved into a global cultural phenomenon with hundreds of millions of readers and a mature commercial ecosystem [19, 31, 37]. However, long-form online novelists, particularly within the Chinese context, face significant challenges in maintaining narrative consistency, coherent character development, and world-building [25] across works spanning millions of words [21]. To meet reader demand and ensure commercial viability, authors must release frequent updates, with each installment adding substantial new content [5, 36]. Moreover, web publication places authors within a more rapid feedback cycle in which readers scrutinize each update, identify inconsistencies, and frequently pressure authors to revise plot elements [22, 39]. In such a consumer-driven ecosystem [19], even minor inconsistencies can erode reader trust and immersion, directly threatening reader retention and commercial sustainability. This marathon-like creative process makes even minor oversights in established details potentially disruptive, imposing substantial cognitive load and interrupting the author's creative flow [9, 29, 30].

Facing such immense creative pressure, authors have an urgent need for highly efficient support tools that can facilitate their writing process. Recent research highlights the potential of Large Language Models (LLMs) for creative text generation [23, 35], and AI-assisted writing tools have emerged to support various stages of writing [12, 33, 38]. For example, Wordcraft integrates AI interactions into drafting, node-based editors such as StoryNode [27] support planning through visual organization, and Dramatron [24] provides a hierarchical workflow from characters to plots. These tools greatly improve authors' writing experience in drafting.

However, current tools do not adequately support long-form online authors in tracking the plot and maintaining narrative consistency. First, most AI writing tools primarily focus on generation or

drafting [3, 13, 23], often overlooking authors' creativity [2, 15, 33], and offering limited support for post-writing review and revision, where narrative inconsistencies are commonly detected. For example, VISAR [40] and Nabokov et al.'s work [4] support outline-driven or nonlinear pre-writing, yet both leave post-writing review unsupported. Although some strategies [20, 26] improve controllable long-text generation, authors still need to manually track large amounts of information, with little tool support for auditing coherence. Second, ensuring consistency in online long-form novel is uniquely difficult because it depends on an author-created world and an evolving narrative history that continuously expands across chapters. While some tools [14, 17] improve revision, they mainly support local text-level rewriting rather than verifying consistency against large-scale narrative memory. Prior work also discusses plot holes and contradictions from different perspectives [1, 6, 28], yet these approaches often require highly structured representations or full-text parsing, limiting their flexibility and scalability for long-form narratives. For lengthy web novels, preserving and analyzing story history faces long-context forgetting challenges [11, 18]. Although researchers augment LLMs with external memory via improved retrieval [8] or knowledge graphs [34], these techniques are not well adapted to fiction writing, where "ground truth" is constructed by the author and consistency violations are context-dependent. Therefore, new designs are needed to better support history-aware narrative memory and consistency auditing in online long-form writing.

Building on this perspective, we ask: How can we design a human-AI collaborative system that externalizes authors' narrative memory and audits inconsistencies during long-form writing? To answer this question, we conducted a formative study with professional web novelists to identify key challenges in maintaining long-form narrative consistency. These insights informed a "Write, Audit, Correct" workflow, which we instantiate in TRACE (Tracking, Retrieving, and Auditing for Coherent Epics), a web-based writing system for serialized online long-form novel. While recent works also explore card-based representations for story organization [10, 32], they primarily focus on summarization and information retrieval. TRACE integrates structured external memory and on-demand auditing into the writing workflow by transforming long-form text into dynamic character cards that track evolving plot elements. We contribute (1) formative insights on consistency challenges in online serialized fiction writing, (2) TRACE, a prototype implementing "Write, Audit, Correct" loop with dynamic structured memory and author-triggered auditing, and (3) an early evaluation with 12 authors suggesting implications for author-controlled writing support tools.

2 Formative Study

Method. To understand the needs of professional web novelists, we conducted a two-stage formative study to identify their key pain points in web novel creation, examine what forms of assistance they value, and explore how much creative freedom should AI tools maintain when assisting authors. First, we analyzed public threads from major Chinese author community forums¹ (Aug 2022–Jul 2025) for broad pattern identification, collecting 2,850 threads and

¹<https://www.qidianclub.com>, <https://write.qq.com>

31,400 replies; after filtering, we retained 820 core threads. Second, to deep dive into authors' needs, we conducted semi-structured interviews with eight professional authors (1.5–30 years experience; $M=11.88$, $SD=10.37$; Table 1), approved by IRB and compensated with a \$20 gift card. Interviews were conducted in person with open-ended questions and probing follow-ups, with participants' consent. All data were anonymized and thematically analyzed; two researchers co-coded 20% of the data to develop a shared codebook and synthesize themes.

Findings. Across our forum analysis and interviews, we identified three recurring challenges in web novel creation. First, authors face a high cognitive load from relying on **fragmented and static external memory** (e.g., character sheets, timelines, and scattered notes across tools). While note-taking itself is common, the burden comes from repeatedly retrieving and cross-checking details during writing, which disrupts attention and breaks momentum. As one forum user noted, *"I have to stop writing every five minutes to check my notes, which breaks my focus."* Authors do not want more note-taking tools, but a **unified and accessible narrative memory** that can be consulted within the writing workflow. Second, authors frequently encounter **long-range consistency risks** as narratives evolve, including contradictions in character traits, relationships, and plot logic. These issues are often discovered during later review, forcing time-consuming backtracking through prior chapters and character histories, which can increase their anxiety and uncertainty. As P1 explained, *"as the plot progresses and character interactions become more complex, it is easy to forget or confuse certain character details."* Maintaining coherence is not only a writing challenge but also a sustained **tracking problem that grows with narrative length and complexity**. Third, although some participants have tried AI writing tools, they remain skeptical of AI as a *"ghostwriter"* and instead **prefer intent-preserving support**, particularly for tedious but cognitively demanding consistency checking. Participants expressed concerns that AI-generated text may reduce writing quality or harm reader trust. They also critiqued existing tools as providing a *"chaotic patchwork"* of suggestions, and instead hoped AI could support tasks such as expanding their own ideas, managing long novels with large casts, and assisting with syntax or grammar checking. Four authors envisioned a tool that preserves their creative intent like a *"guardian"*, reflecting a strong preference for AI tools to **keep the core creativity human-driven**.

Table 1: Participants in Formative Study. (Exp.: Year of Experience; BG: Background; FT: Full-time; PT: Part-time; AI: Prior AI Tool Use.)

ID	Occupation	Exp.	BG	Genre	AI
P1	Screenwriter	30	FT	Rural Fiction, Documentary	Yes
P2	Screenwriter	2	PT	Fan Fiction, Romance	No
P3	Playwright	5	FT	Theater / Dramatic Writing	No
P4	Screenwriter	3	FT	Romance / Female-Oriented	No
P5	Screenwriter	7	FT	Young Adult / Campus	No
P6	Editor	10	FT	Multiple Genres	Yes
P7	Screenwriter	13	PT	Fantasy	Yes
P8	Writer	25	FT	Literary Fiction, Essays	No

3 System Design

Design Goals. Based on these findings, we derive three design goals for TRACE: **(DG1) Externalized narrative memory**, which consolidates key story elements into a dynamic and structured knowledge base to support quick lookup and reduce authors' memory burden; **(DG2) Author-triggered, non-intrusive auditing**, which allows writers to invoke consistency checks only when needed and keeps auditing closely integrated with the writing workflow; and **(DG3) Guardian, not ghostwriter**, which provides suggestive and controllable feedback to help authors verify and maintain consistency while preserving their creative ownership and decision-making.

User Interface. To support web novelists' fragmented, cross-device writing, we built TRACE as a web-based tool. The interface adopts a three-column layout (Figure 2): chapter navigation (left), immersive writing (center), and an assistance sidebar (right) with character cards and conflict detection. The editor supports saving drafts and triggering on-demand audits; detected issues are highlighted inline to keep attention on the text. Character cards provide editable, structured plot memory for quick lookup, while the conflict panel lists flagged issues with brief explanations and actions to acknowledge or dismiss.

System Architecture. TRACE is a retrieval-augmented pipeline for novel-length manuscripts with incremental memory updates. It consists of two backend components: **(1) Dynamic Knowledge Base.** Based on our formative study, authors often struggle with fragmented memory about their character setting and plots detail. To mitigate this, we construct a dynamic knowledge base that structures and updates narrative knowledge continuously (DG1). TRACE extracted narrative facts as *Character Cards*. Each card is a character-centered narrative memory, and separates invariant attributes (e.g., name, identity) from evolving attributes (e.g., relationships, key events), modeling the story as each character's evolving personal history. Manuscripts are chunked with overlap (1500 words with 200 overlap), embedded (text-embedding-3-small), and indexed in a vector database (e.g., ChromaDB) for semantic retrieval. Upon saving, TRACE extracts schema-based updates, logs them chronologically with metadata, and matches them to existing cards via invariant attributes. High-confidence matches (similarity ≥ 0.9) are appended as incremental, time-ordered updates to the corresponding cards, while uncertain ones are staged for author confirmation. **(2) Consistency Auditing Engine.** Given a target text span, the engine performs on-demand consistency auditing (DG2). It indexes the manuscript at the sentence level and retrieves relevant Character Cards via vector search. The retrieved cards are then compared against newly written sentences to detect potential violations of established narrative constraints using few-shot prompting. We focus on two predefined categories: **Character Conflict** (identity, trait, event, ability) and **Plot Hole** (relation, chronology, setting, meta). Structured information extraction and consistency checking are performed by LLM (DeepSeek-V3 [7]). The model returns structured outputs for each detected issue, including the conflict type, rationale, supporting card references, and sentence-level indices in the draft for UI highlighting.

Human-in-the-Loop Workflow. Figure 1 illustrates TRACE's author-in-the-loop workflow (DG3). During *Write*, authors draft normally while TRACE incrementally updates the Character Cards




Figure 2: The user interface of TRACE. (1) The Three-Panel Layout UI Structure, with a central Writing Space, the Chapter List (Left) and the dynamic support Sidebar (Right). (2) Character Card: This area in the Sidebar displays automatically extracted “Character Cards”, which summarize key narrative elements and character attributes. (3) Conflict Detection: This area displays the results of consistency checks, and the Conflict Detection panel provides detailed context.

from saved text. Uncertain updates are surfaced for authors to confirm or merge, keeping memory editable and author-grounded. During *Audit*, authors trigger an on-demand consistency check. TRACE compares the newly written text against retrieved memory and surfaces potential conflicts with evidence-backed alerts. During *Correct*, alerts are presented as **suggestions rather than automatic edits**. Authors can revise the draft, dismiss alerts as intentional twists, or continue writing, keeping creative decisions fully in the author’s hands.

4 User Study

We conducted a controlled user study to answer three questions:

- RQ1: How does TRACE support consistency maintenance and reduce cognitive load in long-form writing?
- RQ2: How do authors interact with TRACE, and how does it impact their workflow?

Table 2: Participants in Our User Study. Recruiting web novelists through the online novelist communities. (Exp.: Year of Experience; BG: Background.)

ID	Exp.	BG	Genre
P1	6	Amateur	Romance
P2	10	Professional	Immortal Heroes
P3	5	Student	Fan Fiction
P5	2	Amateur	Historical Fiction
P4	7	Amateur	Fan Fiction
P6	5	Student	Contemporary
P7	8	Professional	Science Fiction
P8	5	Professional	Fan Fiction
P9	4	Amateur	Contemporary
P10	6	Professional	Science Fiction
P11	9	Professional	Romance
P12	1	Student	Science Fiction

- RQ3: How does TRACE influence user perceived usability, creative agency, and trust?

Participants. We recruited web novel authors from online communities under IRB approval. All participants had experience writing long-form Chinese online literature across diverse genres (0.5-10 years; $M=5.67$, $SD=2.64$). Demographics are summarized in Table 2. Participants received a \$30 gift card for compensation.

Procedure. We conducted a remote within-subjects study using two parallel fixed-continuation tasks that required participants to extend an existing story world in a fan-fiction-style setting. Participants were given a well-known canon event to quickly recover context and asked to continue the story from two perspectives while incorporating predefined inconsistency triggers. To mitigate familiarity and ordering effects, we counterbalanced the task order and condition assignment across participants. Each participant completed both TRACE and a Baseline condition (Microsoft Word + a note-taking app; optional chatbot use). Each two-hour session included a 15-minute tutorial, two 30-minute writing blocks with brief post-task surveys, and a 30-minute semi-structured interview. We designed the questionnaire and open questions informed by prior work [16]. All sessions were recorded with informed consent.

Finding. The study suggests that TRACE supports long-form online writing by making consistency work visible, lightweight, and author-controlled, rather than forcing writers into heavy manual review or unwanted AI generation. For RQ1, participants reported reduced cognitive load through **less context switching** and a natural **write-check rhythm**. Writers could move between drafting, consulting card-based narrative memory, and running audits as needed. P1 noted there was “no need to switch to other applications,” and P10 described the workflow as “much smoother.” TRACE also served as an **idea trigger**: P2 said revisiting character cards helped “generate new ideas,” while P5 found it useful to “clarify the narrative

constraints for character design and dialogue.” Participants further felt it helped them “organize my thoughts” (P9) and stay “cognitively aligned with the narrative” (P2) during revision. For RQ2, participants framed TRACE as an **assistant/proofreader** rather than a co-author, emphasizing that “core creativity remained human-driven” (P3, P4). They valued its ability to “identify areas for refinement and offering feedback” (P1, P5) and “saving time” (P10, P11). However, feedback highlighted a **tension between proactivity and control**: while some wanted more proactive prompts, others preferred optional and context-sensitive support. Repeated flags could also constrain intentional deviations; P1 noted the system “limited my planned character reversals.” For RQ3, TRACE enabled **more relaxed exploration with higher confidence**. Character cards helped writers recall prior plot points and write more freely: P10 said this “allows me to be freer in creating new stories,” and P1 noted it helped “reduce anxiety.” Enjoyment was closely tied to conflict detection with brief explanations, while participants emphasized configurable reminders to avoid unnecessary interruption (P7).

5 Discussion and Conclusion

Our findings suggest several design implications for AI-assisted narrative writing tools that support long-form consistency without overriding author intent. First, consistency support should be author-triggered, enabling writers to invoke audits when needed, rather than receiving continuous automatic corrections. Second, auditing feedback should be adjustable and tolerant, since overly strict checks may misinterpret intentional character development or narrative twists as errors; thus, systems should provide evidence-linked warnings and lightweight controls (e.g., accept/ignore/mark-as-exception). Third, practical deployment must account for NLP limitations in long narratives, including entity disambiguation (e.g., aliases and ambiguous references), long-context constraints, and contradictions that rely on implicit social rules or commonsense reasoning.

Our study has several limitations. First, our participants were limited to Chinese online literature writers. While this community is distinct and influential, it does not represent the full diversity of potential users. Future work should include writers with different languages, genres, and creative backgrounds. Second, our user study relied on controlled tasks, which improved experimental consistency but did not allow participants to work on their own writing projects. Future studies should evaluate TRACE in real-world settings with authentic creative workflows. Finally, our current system focuses on long-form narrative consistency. Future extensions could support broader auditing needs, such as genre-specific fact checking (e.g., historical fiction), grammar and style refinement, world-level constraints checking, and readability revisions. It may also generalize to other long-form creative formats, such as screenplay writing and episodic storytelling.

References

- [1] Kabir Ahuja, Melanie Sclar, and Yulia Tsvetkov. 2025. Finding Flawed Fictions: Evaluating Complex Reasoning in Language Models via Plot Hole Detection. *CoRR* abs/2504.11900 (2025). arXiv:2504.11900 doi:10.48550/ARXIV.2504.11900
- [2] Oloff C. Biermann, Ning F. Ma, and Dongwook Yoon. 2022. From Tool to Companion: Storywriters Want AI Writers to Respect Their Personal Values and Writing Strategies. In *DIS '22: Designing Interactive Systems Conference, Virtual Event, Australia, June 13 - 17, 2022*. ACM, 1209–1227. doi:10.1145/3532106.3533506
- [3] Alex Calderwood, Vivian Qiu, Katy Ilonka Gero, and Lydia B. Chilton. 2020. How Novelists Use Generative Language Models: An Exploratory User Study. In *Joint Proceedings of the Workshops on Human-AI Co-Creation with Generative Models and User-Aware Conversational Agents co-located with 25th International Conference on Intelligent User Interfaces (IUI 2020), Cagliari, Italy, March 17, 2020 (CEUR Workshop Proceedings, Vol. 2848)*. CEUR-WS.org. <https://ceur-ws.org/Vol-2848/HAI-GEN-Paper-3.pdf>
- [4] DASHIEL CARRERA, ZIXIN ZHAO, ASHISH THOMAS, and DANIEL WIGDOR. 2025. Nabokov’s Cards: An LLM-Assisted Co-Writing System to Support Bottom-Up Creative Writing. (2025).
- [5] Jing Chen. 2012. Refashioning Print Literature: Internet Literature in China. *Comparative Literature Studies* 49 (2012), 537 – 546. <https://api.semanticscholar.org/CorpusID:146775927>
- [6] Aron Davids. 2022. *Identifying plot holes in narrative stories by simulating events*. Thesis. University of Twente, Enschede. <https://purl.utwente.nl/essays/91967>
- [7] DeepSeek-AI. 2024. DeepSeek-V3 Technical Report. *ArXiv* abs/2412.19437 (2024). <https://api.semanticscholar.org/CorpusID:275118643>
- [8] Yiming Du, Hongru Wang, Zhengyi Zhao, Bin Liang, Baojun Wang, Wanjuan Zhong, Zezhong Wang, and Kam-Fai Wong. 2024. PerL2QA: A Personal Long-Term Memory Dataset for Memory Classification, Retrieval, and Fusion in Question Answering. *Proceedings of the 10th SIGHAN Workshop on Chinese Language Processing (SIGHAN-10) (2024)*. <https://api.semanticscholar.org/CorpusID:271769814>
- [9] Linda Flower and John R Hayes. 1981. A cognitive process theory of writing. *College Composition & Communication* 32, 4 (1981), 365–387.
- [10] Kexue Fu, Jingfei Huang, Long Ling, Sumin Hong, Yihang Zuo, Ray LC, and Toby Jia jun Li. 2026. Vistoria: A Multimodal System to Support Fictional Story Writing through Instrumental Text-Image Co-Editing. *arXiv:2509.13646 [cs.HC]* <https://arxiv.org/abs/2509.13646>
- [11] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Qianyu Guo, Meng Wang, and Haofen Wang. 2023. Retrieval-Augmented Generation for Large Language Models: A Survey. *CoRR* abs/2312.10997 (2023). arXiv:2312.10997 doi:10.48550/ARXIV.2312.10997
- [12] Al-Mothana M. Gasaymeh, Mohammad A. Beirai, and Asma’a A. Abu Qbeita. 2024. University Students’ Insights of Generative Artificial Intelligence (AI) Writing Tools. *Education Sciences* (2024). <https://api.semanticscholar.org/CorpusID:272959957>
- [13] Katy Ilonka Gero and Lydia B Chilton. 2019. How a Stylistic, Machine-Generated Thesaurus Impacts a Writer’s Process. In *Proceedings of the 2019 Conference on Creativity and Cognition*. 597–603.
- [14] Md Naimul Hoque, Tasfia Mashiat, Bhavya Ghai, Cecilia D Shelton, Fanny Chevalier, Kari Kraus, and Niklas Elmquist. 2024. The HaLLMark effect: Supporting provenance and transparent use of large language models in writing with interactive visualization. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*. 1–15.
- [15] Taewook Kim, Hyomin Han, Eytan Adar, Matthew Kay, and John Joon Young Chung. 2024. Authors’ Values and Attitudes Towards AI-bridged Scalable Personalization of Creative Language Arts. In *Proceedings of the CHI Conference on Human Factors in Computing Systems, CHI 2024, Honolulu, HI, USA, May 11-16, 2024*. ACM, 31:1–31:16. doi:10.1145/3613904.3642529
- [16] Yewon Kim, Sung-Ju Lee, and Chris Donahue. 2025. Amuse: Human-AI Collaborative Songwriting with Multimodal Inspirations. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems, CHI 2025, Yokohama, Japan, 26 April 2025 - 1 May 2025*. ACM, 187:1–187:28. doi:10.1145/3706598.3713818
- [17] Philippe Laban, Jesse Vig, Marti Hearst, Caiming Xiong, and Chien-Sheng Wu. 2024. Beyond the chat: Executable and verifiable text-editing with llms. In *Proceedings of the 37th Annual ACM Symposium on User Interface Software and Technology*. 1–23.
- [18] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*. <https://proceedings.neurips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html>
- [19] Zhenling Li. 2023. Cross-culture, translation and post-aesthetics: Chinese online literature in/as world literature in the Internet era. *World Literature Studies* (2023). <https://api.semanticscholar.org/CorpusID:263810538>
- [20] Xun Liang, Hanyu Wang, Yezhaohui Wang, Shichao Song, Jiawei Yang, Simin Niu, Jie Hu, Dan Liu, Shunyu Yao, Feiyu Xiong, and Zhiyu Li. 2024. Controllable Text Generation for Large Language Models: A Survey. *CoRR* abs/2408.12599 (2024). arXiv:2408.12599 doi:10.48550/ARXIV.2408.12599
- [21] Yang Liu, Xuru Zhuo, and Xiaozhu Zhou. 2024. Multifractal analysis of Chinese literary and web novels. *Physica A: Statistical Mechanics and its Applications* 641 (2024), 129749.
- [22] Zhenhua Liu, Mingze Xu, and Yi Tang. 2025. The Impact of Real-Time Reader Interactions on Plot Structure and Pacing in Chinese Web Novels. *Studies in Art and Architecture* (2025). <https://api.semanticscholar.org/CorpusID:27691730>

- [23] Piotr Mirowski, Kory W. Mathewson, Jaylen Pittman, and Richard Evans. 2023. Co-Writing Screenplays and Theatre Scripts with Language Models: Evaluation by Industry Professionals. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems, CHI 2023, Hamburg, Germany, April 23–28, 2023*. ACM, 355:1–355:34. doi:10.1145/3544548.3581225
- [24] Piotr Mirowski, Kory W. Mathewson, Jaylen Pittman, and Richard Evans. 2023. Co-writing screenplays and theatre scripts with language models: Evaluation by industry professionals. In *Proceedings of the 2023 CHI conference on human factors in computing systems*. 1–34.
- [25] Zhanghe Ni. 2020. Xiuzhen (Immortality Cultivation) Fantasy: Science, Religion, and the Novels of Magic/Superstition in Contemporary China. *Religion* 11 (2020), 25. <https://api.semanticscholar.org/CorpusID:210978873>
- [26] Kyeongman Park, Nakyeong Yang, and Kyomin Jung. 2024. LongStory: Coherent, Complete and Length Controlled Long Story Generation. In *Advances in Knowledge Discovery and Data Mining - 28th Pacific-Asia Conference on Knowledge Discovery and Data Mining, PAKDD 2024, Taipei, Taiwan, May 7–10, 2024, Proceedings, Part II (Lecture Notes in Computer Science, Vol. 14646)*. Springer, 184–196. doi:10.1007/978-981-97-2253-2_15
- [27] Hua Xuan Qin, Guangzhi Zhu, Mingming Fan, and Pan Hui. 2025. Toward Personalizable AI Node Graph Creative Writing Support: Insights on Preferences for Generative AI Features and Information Presentation Across Story Writing Processes. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems, CHI 2025, Yokohama Japan, 26 April 2025– 1 May 2025*. ACM, 897:1–897:30. doi:10.1145/3706598.3713569
- [28] Marie-Laure Ryan. 2009. Cheap Plot Tricks, Plot Holes, and Narrative Design. *Narrative* 17 (2009), 56 – 75. <https://api.semanticscholar.org/CorpusID:49387599>
- [29] John Sweller. 2011. Cognitive load theory. In *Psychology of learning and motivation*. Vol. 55. Elsevier, 37–76.
- [30] Xiaoli Tian and Michael Adorjan. 2016. Fandom and coercive empowerment: the commissioned production of Chinese online literature. *Media, Culture & Society* 38 (2016), 881 – 900. <https://api.semanticscholar.org/CorpusID:147276107>
- [31] Michael S. C. Tse and Maleen Z. Gong. 2012. Online Communities and Commercialization of Chinese Internet Literature. *Journal of Internet Commerce* 11 (2012), 100 – 116. <https://api.semanticscholar.org/CorpusID:155034460>
- [32] Xiyuan Wang, Yi-Fan Cao, Junjie Xiong, Sizhe Chen, Wenxuan Li, Junjie Zhang, and Quan Li. 2025. ClueCart: Supporting Game Story Interpretation and Narrative Inference from Fragmented Clues. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems, CHI 2025, Yokohama Japan, 26 April 2025– 1 May 2025*. ACM, 410:1–410:26. doi:10.1145/3706598.3713381
- [33] Utami Widiati, Diyenti Rusdin, Inda Indrawati, Marzuki, and Nadaraj Govender. 2023. The impact of AI writing tools on the content and organization of students' writing: EFL teachers' perspective. *Cogent Education* 10 (2023). <https://api.semanticscholar.org/CorpusID:260031399>
- [34] Yike Wu, Nan Hu, Sheng Bi, Guilin Qi, Jie Ren, Anhuan Xie, and Wei Song. 2023. Retrieve-Rewrite-Answer: A KG-to-Text Enhanced LLMs Framework for Knowledge Graph Question Answering. *CoRR abs/2309.11206* (2023). arXiv:2309.11206 doi:10.48550/ARXIV.2309.11206
- [35] Haotian Xia, Hao Peng, Yunjia Qi, Xiaozhi Wang, Bin Xu, Lei Hou, and Juanzi Li. 2025. StoryWriter: A Multi-Agent Framework for Long Story Generation. *CoRR abs/2506.16445* (2025). arXiv:2506.16445 doi:10.48550/ARXIV.2506.16445
- [36] Jinghong Xiao, Dingjia Chen, Qiao Tang, Luo Suo, Xingjie Zhou, Wenjing Wang, Wei Zheng, and Man Zhou. 2022. 2020 Report on the development of Chinese Internet literature. *New Techno Humanities* 2, 1 (2022), 1–12.
- [37] Guobin Yang. 2010. Chinese Internet Literature and the Changing Field of Print Culture. <https://api.semanticscholar.org/CorpusID:192965618>
- [38] Catherine Yeh, Gonzalo A. Ramos, Rachel Ng, Andy Huntington, and Richard Banks. 2024. GhostWriter: Augmenting Collaborative Human-AI Writing Experiences Through Personalization and Agency. *CoRR abs/2402.08855* (2024). arXiv:2402.08855 doi:10.48550/ARXIV.2402.08855
- [39] Ze Yu and Federico Pianzola. 2024. Across the Pages: A Comparative Study of Reader Response to Web Novels in Chinese and English on Qidian and WebNovel. In *Proceedings of the Computational Humanities Research Conference 2024, Aarhus, Denmark, December 4–6, 2024 (CEUR Workshop Proceedings, Vol. 3834)*. CEUR-WS.org, 322–333. <https://ceur-ws.org/Vol-3834/paper63.pdf>
- [40] Zheng Zhang, Jie Gao, Ranjodh Singh Dhaliwal, and Toby Jia-Jun Li. 2023. Visar: A human-ai argumentative writing assistant with visual programming and rapid draft prototyping. In *Proceedings of the 36th annual ACM symposium on user interface software and technology*. 1–30.